
Decision Tree, Random Forest, and XGBoost

Jindal K. Shah (jindal.shah@okstate.edu)

School of Chemical Engineering
Oklahoma State University

9th iCoMSE Workshop: Machine Learning for Molecules



Jindal K. Shah


Topics

- Decision Tree
- Bagging (bootstrap + aggregating)
- Random forests
- Hands-on Exercise
- Gradient Boosting and extreme Gradient Boosting (XGBoost)




Decision Tree

- Supervised learning method
- Regression/classification
- Non-linear model
- If...then...else...
- Non-parametric model
- Piecewise continuous



Decision Tree



**Depth of the tree
= maximum
number of
branches to
reach a leaf**

Objective function

- The objective is to minimize the residual sum of squares

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

- Here J represents the number of regions the feature space is partitioned into. The prediction in each of the regions is given by the average response.


$$\hat{y}_{R_j} = \frac{\sum_{i \in R_j} y_i}{N_{R_j}}$$

Partitioning the Feature Space

- Top-down greedy approach known as *recursive binary splitting*:
 - Begins at the top of the tree and successively splits the feature space
 - Greedy because the split at a particular step minimizes the RSS at that step rather than splitting in such a way to achieve a better tree in a future step
- Consider a split over a feature j and the corresponding threshold value s , which divides the data such that

$$R_1(j, s) = \{X | X_j < s\}$$

$$R_2(j, s) = \{X | X_j \geq s\}$$




Selection of feature and its threshold

- Minimize the RSS

$$\sum_{i \in R_1(j,s)} (y_i - \hat{y}_{R_1})^2 + \sum_{i \in R_2(j,s)} (y_i - \hat{y}_{R_2})^2$$

$i = 1, 3, 4$ $i = 2, 5$

$\frac{y_1 + y_3 + y_4}{3}$ $\frac{y_2 + y_5}{2}$



(Dis)Advantages of Decision-Tree Models

- Advantages

- Ease of interpretation
- Graphical representation and understanding by a non-expert
- Scaling of features is not required

- Disadvantages

- Accuracy is usually lower than other regression-based approaches
- Small changes in the data can greatly impact the tree structure
- As outputs are only piecewise continuous, multiple inputs can yield identical results.

Overcoming disadvantages

- Bagging (bootstrap + aggregating)
 - Using multiple decision-tree models
- Random forests

Bootstrap sampling

- Using the same data set, create multiple data sets by randomly drawing samples with replacement



Original Data



Bootstrapped Sample



Bootstrapped Sample

Aggregating

- For each of the bootstrapped data set i , develop a decision-tree model and predict a response $f_i(x)$
- Average each of the responses to obtain the response due to bagging.

$$f_{\text{bag}}(x) = \frac{\sum_i f_i(x)}{B}$$

Random Forests

- Multiple decision-tree models
- Bootstrapped data set
- Randomly selected subset of features at every split
- Achieves decorrelation of trees
- Hyperparameters:
 - Number of trees
 - Number of features to select at every split
 - Minimum number of samples required at an internal node
 - Minimum number of samples required at a leaf node

Hands-on Exercise

- RandomForest.ipynb
- Dataset: Surface tension of deep eutectic solvents



Gradient Boosting

- Borrow concept from RF but build trees sequentially
- Idea is to fit to residuals from the previous prediction
- Consider the following dataset

$$\{x_1, y_1\}, \{x_1, y_1\}, \{x_1, y_1\}, \dots, \{x_n, y_n\}$$

- In the first step, response for each step is predicted to be the average response

$$y_i^0 = \frac{\sum y_i}{N}$$

- Residual for each of the data point is computed as

$$r_i^0 = y_i - y_i^0$$

Jindal K. Shah



Gradient Boosting

- A decision-tree is obtained for the residuals, which provides an estimate of the residual for the i^{th} datapoint, say \hat{r}_i^1

- New prediction = old prediction + learning parameter * residual prediction

$$\hat{y}_i^1 = y_i^0 + \nu * \hat{r}_i^1$$

- New residual = Output – New prediction $r_i^1 = y_i - \hat{y}_i^1$

- Fit a decision-tree to r_i^1 and update predictions

- As one might imagine, the number of trees becomes a hyperparameter


Extreme Gradient Boosting

- Very similar to the gradient boosting but the split is based on similarity score and gain
- As before, compute residuals: $r_i^0 = y_i - y_i^0$
- Compute similarity score

$$\frac{\sum_i^{N_r} r_i^0}{\text{No. of residuals} + \lambda}$$


Regularization parameter

Splitting a Node in XGBoost



Step through different values of the threshold and features; select the pair that maximizes Gain.

Output of a Leaf



New predictions are obtained in a similar manner as that for the gradient boosting method – Slide 16

Thank you!



Jindal K. Shah

22