


Dimensionality Reduction



Shuwen Yue



Assistant Professor, Cornell University

May 1, 2025

What do we use it for?


- Visualizing structure in high-dimensional data
- Identifying patterns or clusters
- Removing noise or redundant features
- Reconstruction or generation
- PCs are often used as features in ML models

Clustering




Approaches


PCA



t-SNE






UMAP

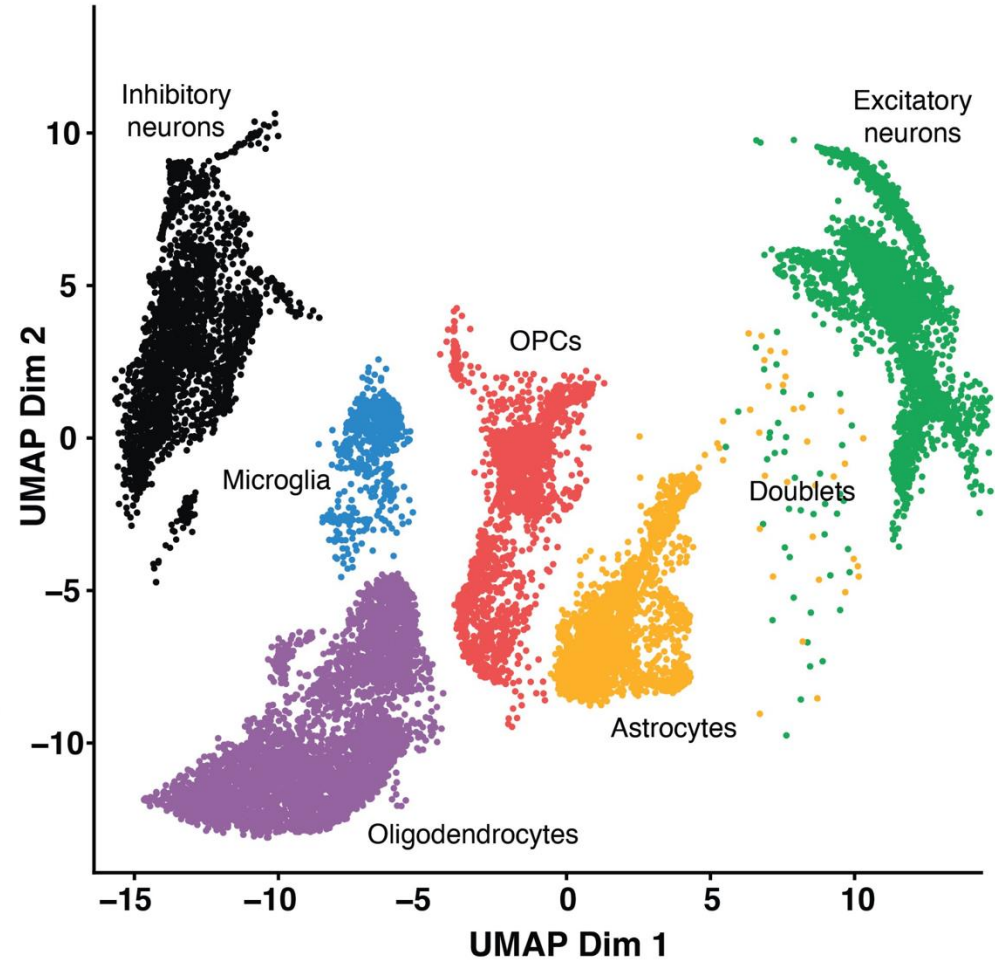
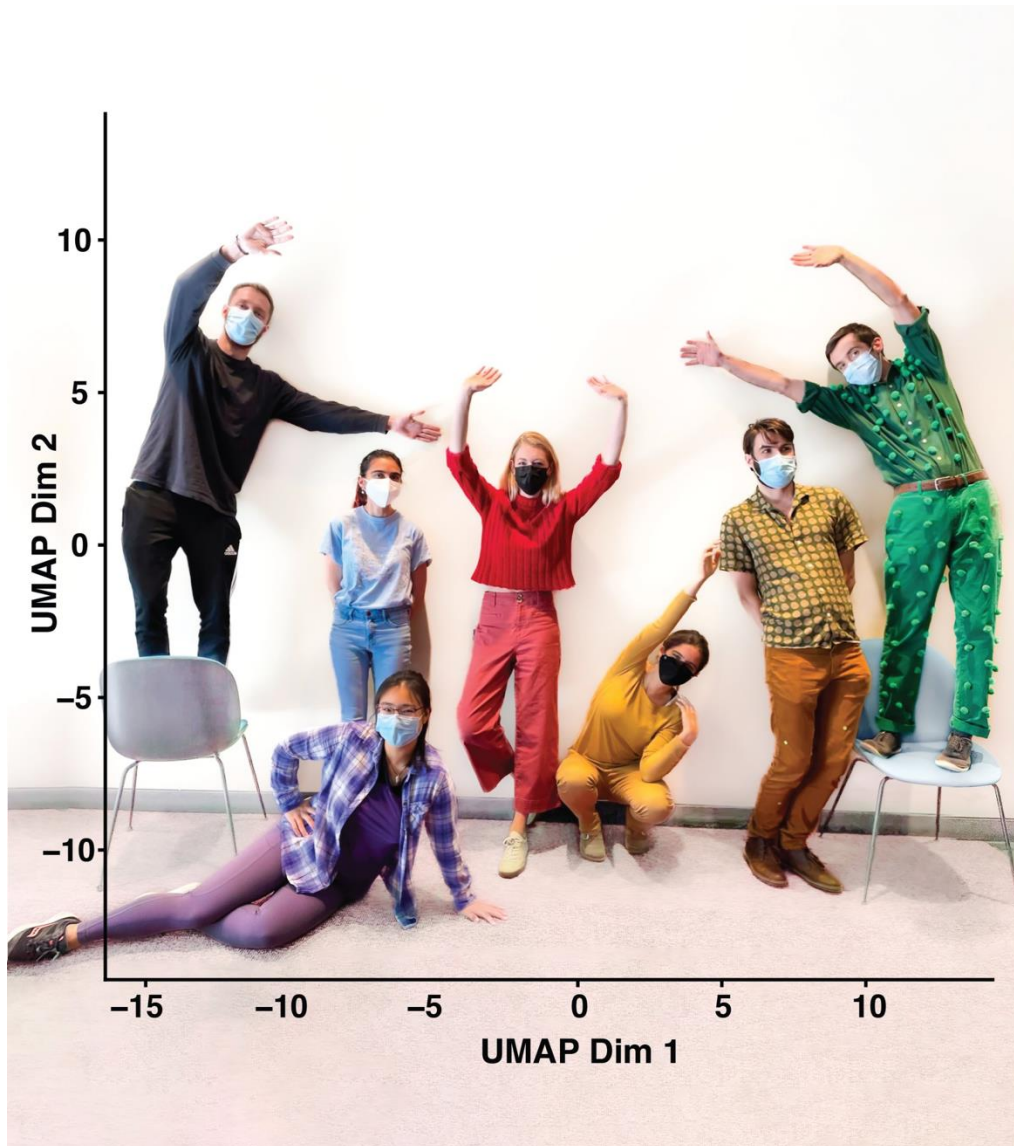


<https://arize.com/blog-course/reduction-of-dimensionality-top-techniques/>

For an example dataset of 2683 aryl bromides from Reaxys, using RDKit fingerprints



Research group Halloween costume idea...




Retaining information – no free lunch!

Distanced preservation – ideally, we want things that are close together in high dimensional space to also be close together in low dimensional space

PCA


Spearman Correlation for PCA: 0.574



PCA distance

t-SNE


Spearman Correlation for tSNE: 0.488



tSNE distance

UMAP

Spearman Correlation for UMAP: 0.607

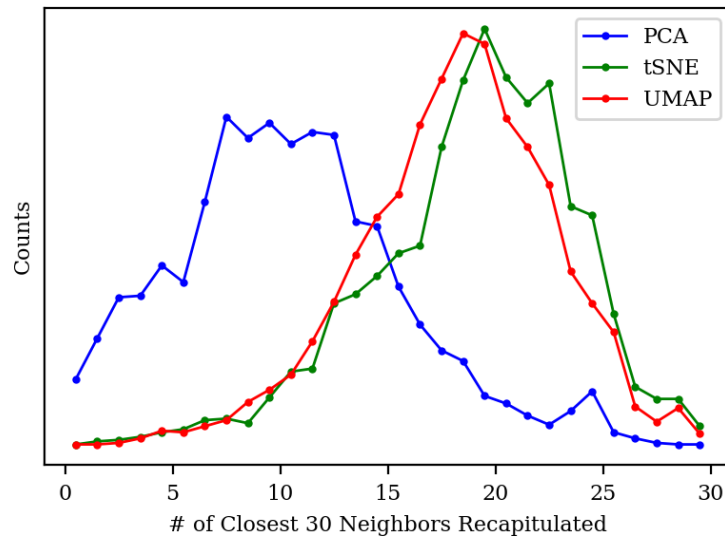


UMAP distance


Retaining information – no free lunch!

While dimensionality reduction can help uncover patterns and importance metric, there is still a fundamental trade off and cost in information loss


30 nearest neighbors



X nearest neighbors




Increasing dimensions



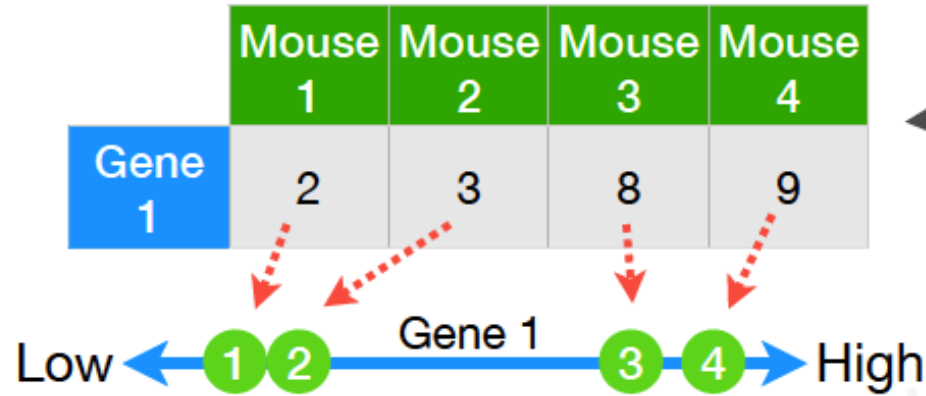
Principle Component Analysis (PCA)

- PCA transforms data linearly into new properties that are not correlated with each other
- Based on Single Value Decomposition (SVD)




The Problem

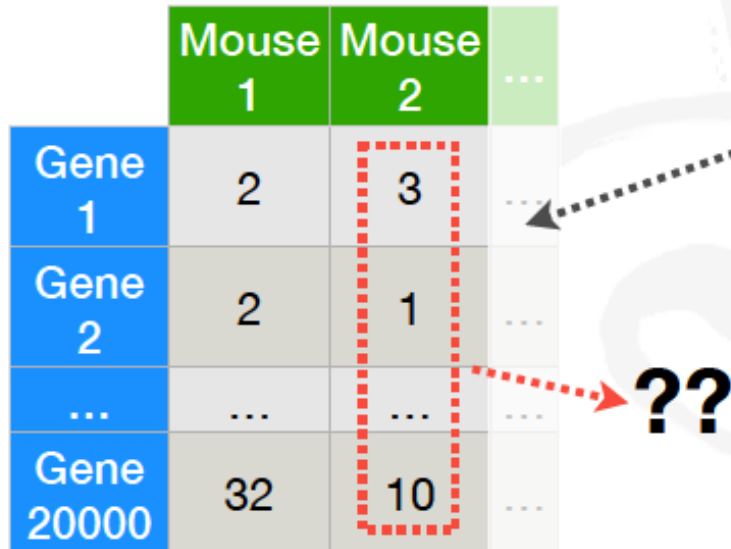
When we have low-dimensional data, like one or two measurements per subject (in this case subject = mouse)...



...graphing each subject is easy...



...but when we have high-dimensional data, like measuring 20,000 genes per mouse, graphing each mouse is not so easy.




PCA, Step-by-Step


First we will demonstrate the concepts with 2-Dimensional data (2 measurements, Gene 1 and Gene 2, per subject).

The data and related graph:

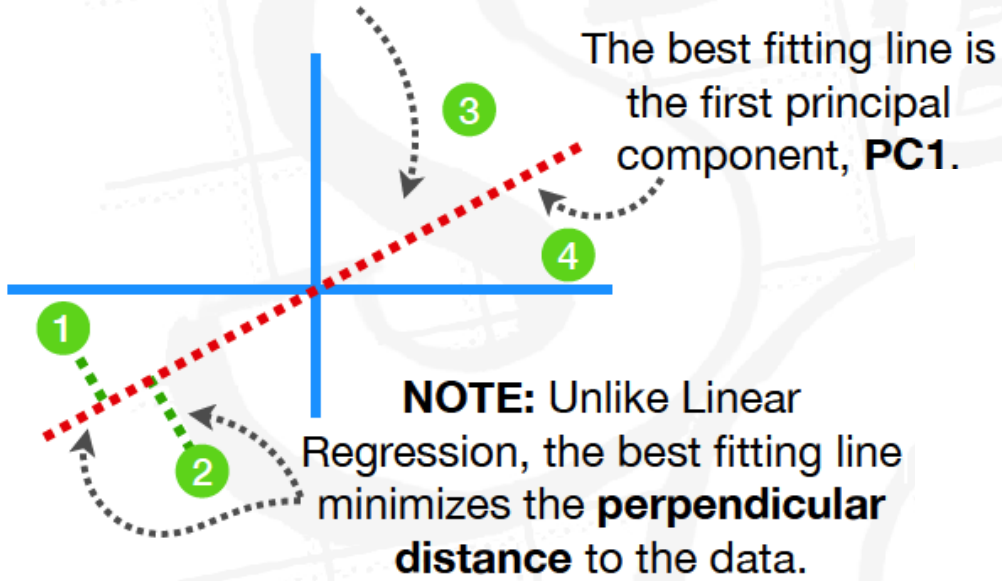
	Mouse 1	Mouse 2	Mouse 3	Mouse 4
Gene 1	2	3	8	9
Gene 2	4	1	7	6




Step 1: Center the Data...




Step 2: Fit a line to the data that goes through the origin...




Due to the **Pythagorean Theorem**, minimizing the perpendicular distance...



The **PC** maximizes the variation between where the perpendicular lines intersect and the origin.






NOTE: The magnitude of Gene 1's **Loading Score** is > the magnitude of Gene 2's. This tells us that Gene 1 is responsible for more variation along **PC1**.

Step 3: Fit a line to the data that is perpendicular to PC1...


The line perpendicular to **PC1** is second principal component, **PC2**.



NOTE: Because the original data only has 2-Dimensions, there are only 2 principal components.

In general, the number of PCs is determined by whichever is smaller, the number of samples (samples = mice in this example) or the number of variables (variables = genes in this example)


A unit vector in the direction of **PC2** is called a **Singular Vector** or **Eigenvector** for **PC2**.




Gene 1	-0.6
Gene 2	0.8

The magnitudes of **Loading Scores** (the coordinates for the unit vector) show that Gene 2 is responsible for more variation along **PC2** than Gene 1.


Step 4: Rotate the data and PCs...



The **light blue lines** are the original axes for Gene 1 and Gene 2.



This is the final **PCA** graph of the data.



The way the data points are spread out along **PC1** tells us that samples (mice) 1 and 2 are more similar to each than they are to samples 3 and 4.


Eigenvalue = $\frac{\text{Sum of Squared Distances along a PC}}{n - 1}$

Variance = $\frac{\text{Sum of Squared Distances along a PC}}{n - 1}$

...where **n** is the number of data points.


If we have 3-dimensions...

	Mouse 1	Mouse 2	Mouse 3	Mouse 4
Gene 1	2	3	8	9
Gene 2	4	1	7	6
Gene 3	8	10	-3	-4




NOTE: The axis for Gene 3 is supposed to represent the 3rd dimension, which is hard to draw on a 2-D piece of paper. Just try to imagine it sticking out of the page.


In 3-D, **PC1** is still the best fitting line that goes through the origin...



...and it accounts for the largest **Eigenvalue**, the sum of the squared distances.

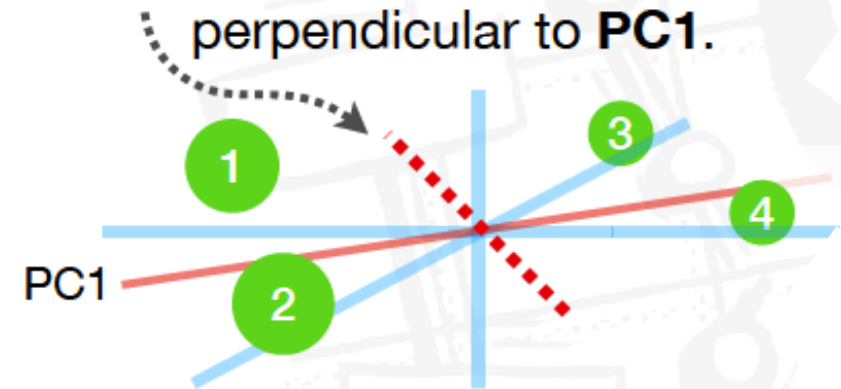


...however, now the **Eigenvector**, the unit vector in the direction of the **PC**, has 3 coordinates.





Because Gene 1's coordinate (**Loading Score**) has the largest magnitude, Gene 1 plays the largest role in the direction of **PC1**.

PC2 is the next best fitting line, given that it goes through the origin and is perpendicular to **PC1**.



PC3 is the next best fitting line, given that it goes through the origin and is perpendicular to **PC1** and **PC2**.







A **Scree Plot** shows the **Eigenvalues**, or if divided by $n-1$, the variances, for the **PCs**.

Scree Plots help us evaluate how many **PCs** we need to accurately represent the original data.

To convert the 3-D graph into a 2-D PCA plot...



Step 1: Look at the Scree Plot: The **Scree Plot** tells us how much variation each **PC** accounts for.




In this case **PC3** accounts for only a very small amount of the variation and it makes sense to exclude it.

In contrast, if this were the **Scree Plot**, we might hesitate to exclude **PC3** since it accounts for almost as much variation as **PC2**




Step 2: Project the data onto PC1 and PC2

Perpendicular lines from the data to **PC1** show the projection onto **PC1**.

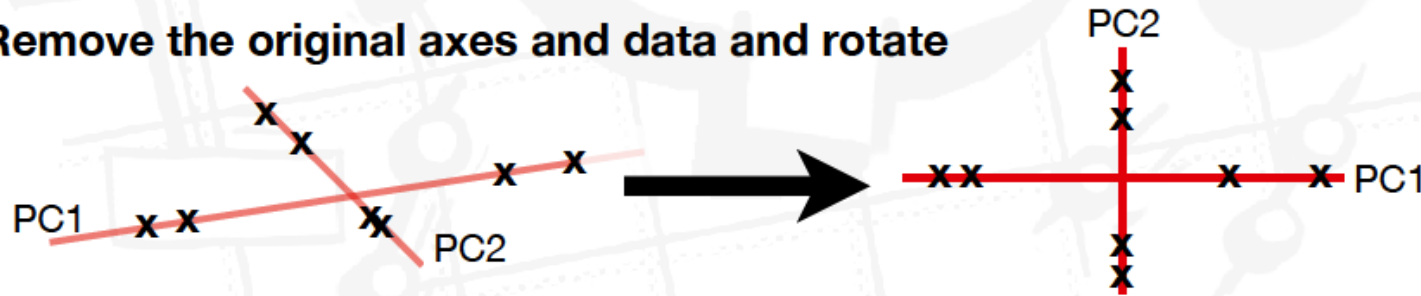


Perpendicular lines from the data to **PC2** show the projection onto **PC2**.




The **x**'s show the projection of the data onto the **PCs**.

Step 3: Remove the original axes and data and rotate




Step 4: Add the data back


These two **x**'s are
for data point 1.



...so point 1 goes
here.



Likewise, the remaining
points are drawn



Colab notebook